

Unlocking Customer Value Through Behavioural Segmentation

RFM Analysis — Customer Intelligence Report

Analysis across 100 customers | 8 of 11 behavioural segments present | 3-year CLV horizon

Executive Summary

Three findings define the commercial opportunity in this customer base

01

Champions & Loyal represent 15% of customers but deliver 40% of total portfolio CLV

15 customers — avg CLV \$1,412 (Champions) and \$832 (Loyal). Protecting and deepening these relationships is the highest-return priority.

02

40% of the base shows lapsing signals — representing recoverable CLV at risk

40 customers across Hibernating, Lost, At Risk, and Can't Lose Them segments. Targeted win-back programmes could recapture a meaningful share of this value before permanent churn.

03

40% of customers are in early-stage growth segments with conversion potential

40 customers across New Customer, Promising, and Potential Loyalist. Structured onboarding and loyalty nudges can migrate these cohorts toward higher-value segments.

Methodology

RFM scoring translates transactional data into actionable behavioural profiles

1

R — Recency

Days since last purchase.
Lower recency = higher score (1–5).
Customers scored via percentile rank.

2

F — Frequency

Number of distinct purchases.
Direct ladder: 1 tx → score 1,
5+ tx → score 5.

3

M — Monetary

Total spend to date.
Higher spend = higher score (1–5).
Customers scored via percentile rank.

4

Segmentation

11 segments assigned via
R×F score matrix — from
Champions to Lost.

5

CLV Estimate

$\text{avg_order} \times \text{freq} \times 3 \text{ yrs} \times 20\%$
margin. Consistent 3-year
horizon for comparability.

Configuration

Input file:	rfm_mockup.csv	Customers analysed:	100	CLV lifespan:	3 years
Margin assumption:	20%	Segments:	11	Output file:	rfm_analysis.csv


Segment Landscape

The 100-customer base spans 11 segments with a total estimated portfolio CLV of \$41,135


Segment	Customers	Share	Avg CLV	Total CLV	Avg Recency	Strategic Priority
Champions	7	7%	\$1,412	\$9,882	30.0 days	Protect & deepen — highest CLV, highest loyalty
Loyal	8	8%	\$832	\$6,655	62.0 days	Reward & upsell — consistent, high-value base
Potential Loyalist	15	15%	\$599	\$8,992	67.0 days	Nudge to loyalty — targeted incentives
Need Attention	5	5%	\$512	\$2,562	134.0 days	Monitor
New Customer	10	10%	\$271	\$2,706	88.0 days	Welcome & educate — first impression drives LTV
Promising	15	15%	\$244	\$3,654	147.0 days	Onboard & engage — early-stage, high conversion upside
Hibernating	20	20%	\$192	\$3,842	230.0 days	Re-engage — time-sensitive before permanent churn
Lost	20	20%	\$142	\$2,843	333.0 days	Win-back — selective recovery, low probability

Champions and Loyal customers (15% of base) generate 40% of total portfolio CLV

The classic 80/20 dynamic is present — concentration at the top is both an asset and a risk



Total CLV by Segment



40% of customers show lapsing behaviour — \$6,685 in CLV is at risk of permanent loss

Hibernating and Lost segments require time-sensitive intervention to prevent irreversible churn

Hibernating

20 customers

Avg CLV: \$192

Total CLV at risk: \$3,842

Avg recency: 230.0 days

Last purchase 193–280 days ago. Single-purchase customers.
High reactivation potential with targeted offer.

Lost

20 customers

Avg CLV: \$142

Total CLV at risk: \$2,843


Avg recency: 333.0 days

Last purchase 280–365 days ago. Lowest recency scores.
Selective win-back only — focus on higher CLV profiles.

⚠ Recommended action: Launch a re-engagement sequence within 30 days for Hibernating customers. Prioritise the top quartile by CLV (above \$204). A/B test discount vs. content-led reactivation.

40% of the base are early-stage customers — structured nurture can convert them to high-value segments

New Customer, Promising, and Potential Loyalist cohorts represent the pipeline for future CLV growth



Migration target: move Promising and Potential Loyalist customers to Loyal status within 12 months

What the Numbers Could Not See

Reading what customers wrote, alongside what they did

Customers flagged by language

5

Healthy on every purchase measure

Lifetime value exposed

\$3,502

8.5% of portfolio

Need a person, not a campaign

23

of 100 customers reviewed

Raised a serious problem

21

Something went wrong, unresolved

FLAGGED BY WHAT THEY SAID, NOT WHAT THEY BOUGHT


CUSTOMER	SEGMENT	LIFETIME VALUE	WHAT IS DRIVING IT	WHY RFM MISSED THEM
CUS5907I	Loyal	\$1,051	Service	Still buying on schedule, but sitting on an unresolved failure
CUSCRUYF	Loyal	\$762	Competitor	Still buying, while actively weighing up another supplier
CUSKOIXN	Loyal	\$743	Service	Still buying on schedule, but sitting on an unresolved failure
CUS65KXV	Potential Loyalist	\$652	Service	Still buying on schedule, but sitting on an unresolved failure
CUSOR56F	New Customer	\$294	Service	Still buying on schedule, but sitting on an unresolved failure

Why this is invisible to RFM: the model is built entirely from past purchases. A customer who has decided to leave but has not yet stopped buying looks identical to a loyal one — right up until they go. Their words are the only early warning.


Customer messages in this analysis are illustrative sample text, generated to demonstrate the method. Replace with real support and survey data before acting on named accounts.

How the Model Read Your Customers

Every customer message scored on two separate questions



Each dot is a customer. Across: what their buying history shows. Up: what they told us. The shaded corner holds customers who are still buying normally while saying they are unhappy — the group a behavioural model cannot surface.



The two questions do not move together. An unresolved failure scores near the top on severity and near the bottom on intent to leave: those customers are complaining to us, not walking away. Treating the two as one number hides both.


Customer messages in this analysis are illustrative sample text, generated to demonstrate the method. Replace with real support and survey data before acting on named accounts.

What It Found, and How Sure It Was

The reasons behind the numbers, and where the model declined to commit

What is driving dissatisfaction


Across 100 customer messages, read one at a time.



Each message was sorted into a single dominant reason. Just over a third raised no issue at all, which is itself useful: it separates the genuinely content from the merely quiet.

The model says when it is not sure

Low confidence means several plays were defensible — not that the answer is wrong.



The model reports its own certainty. Where one play was clearly right it says so; where several were defensible it does not pretend otherwise, and those go to an analyst rather than onto a recommendation slide.

Customer messages in this analysis are illustrative sample text, generated to demonstrate the method. Replace with real support and survey data before acting on named accounts.

Strategic Recommendations

Plays selected against current data, highest portfolio value first

DEFEND • CHAMPIONS

Formalise a VIP tier

- Dedicated contact, early access, exclusive pricing tier
- 7 customers • \$9,882 CLV • Account management • cost: High
- Urgency 0.2 of 3 • No urgency
- Selected with 0.64 confidence

GROW • POTENTIAL LOYALIST

Onboard into the loyalty programme

- Structured nurture sequence into the loyalty programme
- 15 customers • \$8,992 CLV • Email campaign • cost: Low
- Urgency 0.5 of 3 • No urgency
- ANALYST REVIEW — confidence 0.45, runner-up was 'vip_programme'

DEFEND • LOYAL

Run a service recovery

- Name the failure, fix it, compensate, manager closes the loop
- 8 customers • \$6,655 CLV • Human — manager • cost: Medium
- Urgency 0.6 of 3 • Worth planning
- ANALYST REVIEW — confidence 0.49, runner-up was 'vip_programme'

RECOVER • HIBERNATING

Time-boxed win-back

- Expiring reactivation offer, with a measured control group
- 20 customers • \$3,842 CLV • Email campaign • cost: Medium
- Urgency 1.9 of 3 • Act this month
- Selected with 1.00 confidence

Next Steps

Proposed 90-day roadmap to operationalise findings

Days 1-30 Foundation

- Validate RFM scores against CRM and transaction data
- Socialise segmentation with Marketing, CRM, and Analytics teams
- Define success metrics and baseline KPIs per segment
- Brief CRM team on Champions & Loyal VIP programme design

Days 31-60 Activation

- Launch Hibernating win-back email sequence (test vs. control)
- Deploy second-purchase incentive to New Customer cohort
- Onboard Promising customers into loyalty programme
- Set up monthly RFM refresh cadence in data pipeline

Days 61-90 Optimise

- Review win-back programme results — iterate messaging and offer
- Track segment migration rates vs. targets
- Present first monthly CLV dashboard to leadership
- Define Phase 2 scope: predictive churn modelling, CLV uplift testing

Appendix — Full Segment Data

Source: rfm_analysis.csv | 100 customers | Analysis date: March 2026

Segment	# Customers	% Base	Avg CLV	Total CLV	Avg Recency	Avg Frequency	Avg Monetary
Champions	7	7%	\$1,412	\$9,882	30.0 days	4.6x	\$3,396
Loyal	8	8%	\$832	\$6,655	62.0 days	3x	\$2,144
Potential Loyalist	15	15%	\$599	\$8,992	67.0 days	2x	\$1,443
Need Attention	5	5%	\$512	\$2,562	134.0 days	2x	\$1,487
New Customer	10	10%	\$271	\$2,706	88.0 days	1x	\$693
Promising	15	15%	\$244	\$3,654	147.0 days	1x	\$731
Hibernating	20	20%	\$192	\$3,842	230.0 days	1x	\$725
Lost	20	20%	\$142	\$2,843	333.0 days	1x	\$710

How This Was Built

Reading customer language without training a model

WHAT WE ASKED

Five plain questions about each customer's most recent message

- Are they signalling they want to leave?
- How serious is the problem they describe?
- What is driving it — service, price, the product, a competitor, or something on their side?
- Could we still keep them if we acted this fortnight?
- Does this need a person, or will a campaign do?
- Answers come back as numbers, not paragraphs, so ordinary code can use them directly.

WHY NOT A TRAINED MODEL

The usual approach did not fit the problem we actually had

- A trained classifier learns from examples. It needs thousands of past messages, each labelled by hand by someone who already knew the answer. We have 100 customers and no labels at all.
- Every new question would mean another labelled set and another training run. Here, adding a question took minutes.
- Our first scoring rule was wrong — it treated a furious customer as safe. Fixing it cost nothing, because the answers were already saved and we simply re-scored them. Retraining a model would have meant starting again.

WHAT STAYED ORDINARY CODE

The model never touches a single calculation

- Every number is plain arithmetic: the RFM scores, lifetime value, the thresholds, the weights, who makes the list.
- The model judges meaning. Code owns every calculation and every decision that follows from it.
- This is the step before machine learning, not a replacement for it. Once enough of these judgments are paired with real outcomes, they become inputs to a conventional model.

The model reports how sure it is.

2 of 8 recommended plays came back below our confidence bar. Those are marked for analyst review rather than presented as settled. A model that says "I am not sure" is more useful than one that guesses confidently — and it is the difference between a slide you can defend and one you cannot.

Customer messages in this analysis are illustrative sample text, generated to demonstrate the method. Replace with real support and survey data before acting on named accounts.